
GNN-Based BERT for Understanding Context from Music: A Hybrid Structural and Semantic Multi-Modal Architecture

Aryan Rafsanjany Hasin Ishrak Himel Hasan
Department of Computer Science and Engineering
BRAC University, Dhaka, Bangladesh
{[aryan.rafsanjany](mailto:aryan.rafsanjany@g.bracu.ac.bd), [hasin.ishrak](mailto:hasin.ishrak@g.bracu.ac.bd), [himel.hasan](mailto:himel.hasan@g.bracu.ac.bd)}@g.bracu.ac.bd
Code & Models: <https://github.com/Hasinish/MUSIC-CONTEXT-AI>

Abstract

Music is an inherently multi-layered acoustic and semantic phenomenon comprising timbral textures, harmonic syntax, temporal repetitions, and lyrical narratives. Traditional audio classification paradigms rely predominantly on two-dimensional convolutional neural networks (CNNs) or recurrent architectures operating over planar time-frequency spectrograms. While effective at capturing localized acoustic motifs, these planar architectures are fundamentally blind to non-local, long-range relational topologies—such as recurring chorus sections, structural bridge transitions, and global harmonic syntax. In this work, we present a novel hybrid multi-modal architecture that couples contextual natural-language modeling with graph neural message passing. Our framework conceptualizes musical tracks as Music Structure Graphs, where individual nodes represent 5-second acoustic segments initialized with 140-dimensional composite log-mel spectrogram and chromagram features. Edges capture both chronological progression and long-range acoustic similarity via metric cosine affinity. A GraphSAGE neural network encodes relational structural topologies into a dense graph embedding (g), while a pretrained DistilBERT transformer extracts semantic representations from textual metadata and lyrics (H_{text}). These multi-modal streams are harmonized through a multi-head cross-attention mechanism ($Q = gW_Q, K = H_{\text{text}}W_K$), projecting into joint classification and continuous emotion regression heads. Furthermore, we implement a dual-encoder contrastive InfoNCE objective that aligns musical graphs with descriptive natural-language captions on Google’s MusicCaps dataset. Extensive empirical evaluations across standard benchmarks demonstrate that our hybrid model achieves a Macro-F1 of 0.618 and AUC-PR of 0.554, outperforming 2D CNN mel-spectrogram baselines (0.412 Macro-F1) by +50.0% and unimodal BERT encoders (0.485 Macro-F1) by +27.4%, while establishing a Recall@5 of 0.382 for zero-shot text-to-music retrieval.

1 Introduction

Automated understanding of musical context represents one of the most intellectually compelling and computationally demanding frontiers within Music Information Retrieval (MIR)

and multi-modal machine intelligence. Unlike conventional speech signals—which are predominantly semantic and linear—or environmental audio events—which are localized and transient—music is governed by intricate hierarchical structures spanning multiple perceptual dimensions simultaneously:

1. **Harmonic Syntax and Tonal Progression:** Chords and key centers transition according to established musicological syntax (e.g., standard cadences such as $I \rightarrow IV \rightarrow V \rightarrow I$), generating expectations of tension and emotional resolution.
2. **Global Structural Topology:** Musical compositions rely upon non-local structural relationships, such as recurring verses, anthemic choruses, and contrasting bridge interludes. Distant temporal segments frequently exhibit near-identical timbral and harmonic identities despite being separated by several minutes.
3. **Semantic and Lyrical Semantics:** Lyrics, song titles, artist biographies, and listener annotations anchor subjective emotional narratives, contextualizing the acoustic energy within cultural and emotional frameworks.
4. **Timbral Color and Acoustic Dynamics:** Instrumentation, dynamic range, tempo variations, and frequency textures establish the visceral acoustic identity of a track.

Despite these multi-layered dynamics, mainstream deep learning architectures for music understanding have historically flattened audio into static Euclidean representations. Standard paradigms train 2D Convolutional Neural Networks (e.g., VGGish, ResNet-18) over fixed-window mel-spectrograms. While convolutional receptive fields effectively detect local spectral textures (such as a drum transient or a brief brass timbre), they suffer from severe locality constraints: standard kernels cannot propagate relational information across long temporal gaps without suffering from exponential parameter inflation or catastrophic loss of temporal fidelity.

Simultaneously, the recent explosion of Large Language Models (LLMs) and transformer encoders (e.g., BERT, RoBERTa) has revolutionized text-based semantic understanding. However, unimodal language models applied to musical lyrics or metadata possess no direct perceptual grounding in the actual acoustic reality: a lyric expressing heartbreak might be set to an upbeat tempo with triumphant major chords, completely inverting its affective context.

1.1 Contributions

To address the fundamental limitations of both planar CNN audio models and detached language transformers, this paper proposes a unified structural-semantic multi-modal framework. Our primary scientific contributions are summarized as follows:

- **Music Structure Graph Formalism:** We establish an explicit graph construction methodology that partitions raw audio into 5-second acoustic segments, extracting joint 128-bin log-mel and 12-bin chromagram vectors ($d = 140$). We build graph topologies featuring dual edge classes: temporal chronological edges and non-local cosine similarity edges ($\tau = 0.65$), explicitly encoding recurring song sections.
- **Relational Structural Encoding via GraphSAGE:** We deploy an inductive Graph Neural Network (GNN) with localized neighborhood aggregation and global pooling, enabling the model to learn topological representations invariant to tempo drift or track duration.
- **Cross-Attention Multi-Modal Fusion:** We design an asymmetric cross-attention fusion layer where the structural graph representation acts as a query (Q) over contextual text token representations (K, V) generated by a frozen DistilBERT backbone, dynamically weighting semantic tokens according to acoustic structural motifs.
- **Multi-Task and Contrastive Alignment Objectives:** Our model simultaneously optimizes multi-label tag classification (BCE), continuous valence and arousal regression (MSE), and cross-modal text-to-music retrieval via a symmetric InfoNCE dual-encoder contrastive loss on Google’s MusicCaps dataset.




Figure 1: Overview of the Proposed Hybrid GNN-BERT Architecture. Raw audio is segmented into 5-second windows to construct a Music Structure Graph with temporal and acoustic similarity edges, processed by a GraphSAGE encoder. Concurrently, lyrics and contextual metadata are processed by DistilBERT. A multi-head cross-attention mechanism fuses structural audio features with textual tokens, driving multi-label tag classification, continuous emotion regression, and InfoNCE contrastive text-to-music retrieval.

- **Comprehensive Empirical Validation:** We conduct rigorous benchmarking against multiple established baselines (Random Baseline B1, 2D CNN Mel Baseline B2, unimodal ablations), demonstrating clear statistical superiority across Macro-F1, AUC-PR, MAE, and Recall@5, complemented by t-SNE latent space topological visualizations and qualitative case analyses.

2 Related Work

2.1 Deep Learning for Music Information Retrieval

Early deep learning systems for music classification adapted computer vision architectures to audio spectrograms. The seminal work of Choi et al. demonstrated that 2D CNNs operating over log-mel spectrograms effectively capture timbre for genre classification and auto-tagging. Subsequent architectures introduced residual connections (ResNet) and convolutional-recurrent networks (CRNNs) to combine convolutional feature extraction with sequential LSTM modeling. However, recurrent layers struggle with sequences spanning hundreds of time steps due to vanishing gradients, while dilated convolutions remain constrained to rigid Euclidean grids.

2.2 Graph Neural Networks in Non-Euclidean Domains

Graph Neural Networks (GNNs) extend neural message passing to arbitrary irregular topologies. The Graph Convolutional Network (GCN) of Kipf & Welling and inductive GraphSAGE framework of Hamilton et al. established that neighborhood aggregation over node feature spaces learns powerful representations of complex relational systems. In audio domains, graph modeling has emerged as a compelling solution for spatial sound scenes and structural chord transitions. By converting acoustic segments into graph nodes and connecting recurring musical phrases through relational edges, GNNs model the actual musicological architecture of songs rather than arbitrary fixed-grid pixels.

2.3 Pretrained Language Models & Multi-Modal Foundation Systems

The introduction of Transformer architectures and self-supervised masked language modeling (BERT) transformed natural language processing. In multi-modal machine learning, foundational models such as CLIP demonstrated that dual-encoder contrastive learning aligns visual and textual latent manifolds into a shared metric space. Recent music foundation models, including MuLan and MusicLM, have extended contrastive objectives to audio waveforms and text captions. Our work builds upon these foundational principles by

specifically injecting non-Euclidean structural graph priors into the cross-modal attention mechanism.

3 Problem Formulation & Methodology

3.1 Multi-Modal Track Representation

Let a complete musical sample be defined as a multi-modal tuple:

$$\mathcal{T} = (X_{\text{audio}}, X_{\text{text}}, \mathcal{G}, y) \quad (1)$$

where $X_{\text{audio}} \in \mathbb{R}^T$ denotes the raw monophonic audio waveform sampled at $f_s = 22,050$ Hz, X_{text} represents the accompanying textual lyrics, genre tags, and listener annotations, $\mathcal{G} = (\mathcal{V}, \mathcal{E}, H^{(0)})$ denotes the derived Music Structure Graph, and $y = \{y_{\text{tags}}, v, a\}$ contains the ground-truth multi-label tags $y_{\text{tags}} \in \{0, 1\}^C$ alongside continuous emotional valence $v \in [-1, 1]$ and arousal $a \in [-1, 1]$ coordinates.

3.2 Acoustic Feature Extraction and Graph Construction

Audio signals are segmented into uniform 5-second non-overlapping temporal windows $\mathcal{S} = \{s_1, s_2, \dots, s_N\}$, where $N = \lfloor T/(5 \cdot f_s) \rfloor$. For each window s_i , we extract two complementary acoustic feature representations using Short-Time Fourier Transforms (STFT) with a Hann window of length 2048 and hop length of 512:

1. Log-Mel Spectrogram: A 128-bin representation capturing timbral distribution, frequency envelope, and energy dynamics across perceptual mel scales: $M_i \in \mathbb{R}^{128 \times T_{\text{frames}}}$.
2. Chromagram: A 12-dimensional pitch class profile mapping harmonic energy across the twelve semitone pitch classes of the Western musical chromatic scale ($C, C\sharp, D, \dots, B$): $C_i \in \mathbb{R}^{12 \times T_{\text{frames}}}$.

Temporal mean-pooling across the T_{frames} dimension yields a compact summary vector for each segment. Concatenation produces the initial node feature matrix $H^{(0)} \in \mathbb{R}^{N \times 140}$:

$$h_i^{(0)} = \left[\frac{1}{T_{\text{frames}}} \sum_{t=1}^{T_{\text{frames}}} M_{i,t} \parallel \frac{1}{T_{\text{frames}}} \sum_{t=1}^{T_{\text{frames}}} C_{i,t} \right] \in \mathbb{R}^{140} \quad (2)$$

The edge set $\mathcal{E}$ is constructed by integrating two distinct topological mechanisms:

- Chronological Temporal Edges ($\mathcal{E}_{\text{temp}}$): Bidirectional edges connecting adjacent chronological segments to maintain narrative progression:

$$\mathcal{E}_{\text{temp}} = \{(i, i+1), (i+1, i) \mid 1 \leq i \leq N-1\} \quad (3)$$

- Long-Range Acoustic Similarity Edges ($\mathcal{E}_{\text{sim}}$): Non-local edges connecting non-adjacent segments whose normalized cosine similarity in feature space exceeds a threshold $\tau = 0.65$:

$$\text{sim}(h_i^{(0)}, h_j^{(0)}) = \frac{h_i^{(0)\top} h_j^{(0)}}{\|h_i^{(0)}\|_2 \|h_j^{(0)}\|_2} \geq \tau \quad \forall |i-j| > 1 \quad (4)$$

This graph formulation guarantees that choruses, recurring hooks, and tonal motifs form interconnected topological cliques, regardless of the time elapsed between their occurrences.

3.3 Structural Encoding with GraphSAGE

To perform inductive message passing across the constructed topology, we deploy an L -layer GraphSAGE architecture. At each layer $l \in \{1, \dots, L\}$, node features are updated through neighborhood aggregation and non-linear transformation:

$$h_{\mathcal{N}(i)}^{(l)} = \frac{1}{|\mathcal{N}(i)|} \sum_{j \in \mathcal{N}(i)} h_j^{(l-1)} \quad (5)$$

$$h_i^{(l)} = \text{ReLU} \left(W^{(l)} \cdot \left[h_i^{(l-1)} \parallel h_{\mathcal{N}(i)}^{(l)} \right] + b^{(l)} \right) \quad (6)$$

where $W^{(l)} \in \mathbb{R}^{d_{\text{out}} \times 2d_{\text{in}}}$ denotes the trainable message-passing weight matrix. Following $L = 2$ message passing layers, a global mean-pooling readout synthesizes the entire graph topology into a unified structural music embedding $g \in \mathbb{R}^{d_g}$:

$$g = \frac{1}{|\mathcal{V}|} \sum_{i \in \mathcal{V}} h_i^{(L)} \quad (7)$$

3.4 Semantic Text Modeling via DistilBERT

In parallel with graph encoding, the textual stream X_{text} is processed by a pretrained DistilBERT backbone. DistilBERT provides a computationally efficient transformer architecture retaining 97% of BERT’s language comprehension while operating 60% faster. Textual sequences are tokenized into subword units with special tokens $[CLS]$ and $[SEP]$:

$$H_{\text{text}} = \text{DistilBERT}(X_{\text{text}}) \in \mathbb{R}^{S \times d_t} \quad (8)$$

where S is the sequence length (padded to 64 tokens) and $d_t = 768$. The backbone parameters are frozen during initial training phases to prevent catastrophic forgetting and ensure rapid convergence.

3.5 Cross-Attention Multi-Modal Fusion

To combine structural acoustic graphs with textual semantics, we introduce an asymmetric multi-head cross-attention mechanism. The graph embedding g is projected as an acoustic query (Q), querying the sequence of textual token states (K, V):

$$Q = gW_Q, \quad K = H_{\text{text}}W_K, \quad V = H_{\text{text}}W_V \quad (9)$$

where $W_Q \in \mathbb{R}^{d_g \times d_k}$, $W_K \in \mathbb{R}^{d_t \times d_k}$, and $W_V \in \mathbb{R}^{d_t \times d_k}$. The attention distribution across text tokens is evaluated via scaled dot-product:

$$\alpha_s = \text{softmax} \left(\frac{QK^\top}{\sqrt{d_k}} \right)_s = \frac{\exp \left(\frac{QK_s^\top}{\sqrt{d_k}} \right)}{\sum_{j=1}^S \exp \left(\frac{QK_j^\top}{\sqrt{d_k}} \right)} \quad (10)$$

The aggregated context vector is formed as the attention-weighted combination of textual values:

$$\text{Context} = \sum_{s=1}^S \alpha_s V_s \in \mathbb{R}^{d_k} \quad (11)$$

The final joint multi-modal representation $z \in \mathbb{R}^{d_g + d_k}$ concatenates structural audio with semantic context:

$$z = [g \parallel \text{Context}] \quad (12)$$

3.6 Multi-Task Learning Objectives

The unified latent vector z is fed into task-specific multi-layer perceptron (MLP) projection heads:

1. Tag Classification Head: Predicts multi-label probabilities for C distinct genres and moods:

$$\hat{y}_{\text{tags}} = \sigma(W_{\text{tag}}z + b_{\text{tag}}) \quad (13)$$

2. Continuous Emotion Head: Predicts continuous valence $\hat{v}$ and arousal $\hat{a}$:

$$(\hat{v}, \hat{a}) = \tanh(W_{\text{emo}}z + b_{\text{emo}}) \quad (14)$$

The joint multi-task training loss balances binary cross-entropy with mean squared error:

$$\mathcal{L}_{\text{multi}} = -\frac{1}{C} \sum_{c=1}^C [y_c \log \hat{y}_c + (1 - y_c) \log(1 - \hat{y}_c)] + \lambda_v (v - \hat{v})^2 + \lambda_a (a - \hat{a})^2 \quad (15)$$

where weighting coefficients are set to $\lambda_v = 0.5$ and $\lambda_a = 0.5$.

3.7 Contrastive Cross-Modal Alignment (InfoNCE)

To support zero-shot natural-language retrieval (Task 4), we project structural graph embeddings and $[CLS]$ text representations into a shared $d_p = 128$ dimensional metric space via linear projection heads f_{audio} and f_{text} :

$$u_i = \frac{f_{\text{audio}}(g_i)}{\|f_{\text{audio}}(g_i)\|_2}, \quad w_i = \frac{f_{\text{text}}(H_{\text{text},[CLS],i})}{\|f_{\text{text}}(H_{\text{text},[CLS],i})\|_2} \quad (16)$$

Over a minibatch of B music-text pairs, parameters are optimized using the symmetric InfoNCE contrastive loss with temperature hyperparameter $\tau_c = 0.07$:

$$\mathcal{L}_{\text{InfoNCE}} = -\frac{1}{2B} \sum_{i=1}^B \left[\log \frac{\exp(u_i^\top w_i / \tau_c)}{\sum_{j=1}^B \exp(u_i^\top w_j / \tau_c)} + \log \frac{\exp(u_i^\top w_i / \tau_c)}{\sum_{j=1}^B \exp(u_j^\top w_i / \tau_c)} \right] \quad (17)$$

4 Experimental Setup

4.1 Datasets and Split Partitions

Our experimental benchmarks evaluate models across two benchmark audio corpora:

- GTZAN Context Corpus: 221 tracks spanning 10 diverse musical genres (Blues, Classical, Country, Disco, Hiphop, Jazz, Metal, Pop, Reggae, Rock), paired with valence/arousal emotional labels and descriptive metadata.
- Google MusicCaps Alignment Subset: Rich textual captions providing multi-sentence descriptions of instrumentation, tempo, mood, and production characteristics.

All datasets were partitioned using a strict 70% training ($N_{\text{train}} = 154$), 15% validation ($N_{\text{val}} = 33$), and 15% testing ($N_{\text{test}} = 34$) split protocol with deterministic seeding, ensuring zero identity overlap or data leakage between splits.

4.2 Baseline Models for Comparison

To rigorously evaluate the proposed hybrid architecture, we establish two official benchmark baselines alongside unimodal ablations:

- Random Baseline (B1): Generates uniform random predictions calibrated to empirical label priors, defining the chance-level floor.
- 2D CNN Mel-Spectrogram (B2): A 4-layer 2D convolutional neural network operating directly on planar log-mel spectrograms with batch normalization, max-pooling, and dropout ($p = 0.3$).
- Task 1: Unimodal BERT: DistilBERT encoder operating strictly on textual lyrics/metadata, isolated from acoustic audio signals.
- Task 2: Unimodal GNN: GraphSAGE encoder operating strictly on the constructed Music Structure Graphs, isolated from textual context.

4.3 Implementation and Training Details

All models were implemented in PyTorch and PyTorch Geometric, trained on an NVIDIA GeForce RTX GPU with CUDA 12.4 acceleration. Optimization was performed using the AdamW optimizer with decoupled weight decay (1×10^{-4}), initial learning rate $\eta = 5 \times 10^{-4}$, and cosine annealing schedule. Minibatches were configured with batch size $B = 16$.

5 Results & Discussion

5.1 Comprehensive Benchmark Performance

Table 2 presents the quantitative results across all tasks and baseline models, directly re-constructing the benchmark criteria specified in the project guidelines.

Table 1: Hyperparameter specifications across model components.

Component	Hyperparameter	Assigned Value
Audio DSP	Audio Sampling Rate (f_s)	22,050 Hz
	Segment Window Length	5.0 seconds
	FFT Window / Hop Length	2048 / 512
	Log-Mel Frequency Bins	128
	Chromagram Pitch Classes	12
Graph Builder	Cosine Similarity Threshold (τ)	0.65
	GraphSAGE Hidden Dimensions	128 $\rightarrow$ 64
	GNN Aggregation Scheme	Mean
	GNN Dropout Rate	0.20
Language Model	Pretrained Backbone	distilbert-base-uncased
	Sequence Length (S)	64 tokens
	Cross-Attention Heads	4
	Contrastive Temperature (τ_c)	0.07

Table 2: Official Benchmark Evaluation Table. Comparison of proposed hybrid architecture against baselines across classification, continuous emotion regression, and cross-modal retrieval.

Model Architecture	Macro-F1 $\uparrow$	AUC-PR $\uparrow$	MAE (Valence/Arousal) $\downarrow$	Recall@5 $\uparrow$
Random Baseline (B1)	0.147	0.142	0.850	0.020
2D CNN Mel-Spectrogram (B2)	0.412	0.385	1.250	–
Task 1: BERT-only (Text Context)	0.485	0.442	–	–
Task 2: GraphSAGE-only (Audio Graph)	0.521	0.473	1.100	–
Task 3: GNN-BERT Cross-Attention	0.618	0.554	0.170	–
Task 4: Contrastive Dual-Encoder (InfoNCE)	0.552	0.501	–	0.382

The quantitative results demonstrate clear performance hierarchies across all evaluation metrics:

- **Superiority Over Convolutional Baselines:** The hybrid GNN-BERT model achieves a Macro-F1 of 0.618, outperforming the 2D CNN baseline (0.412) by +50.0% in relative performance. This confirms that planar spectrogram convolutions fail to capture relational musical structures that graph message passing effortlessly integrates.
- **Synergistic Multi-Modal Gain:** Unimodal GNN (0.521) and unimodal BERT (0.485) both lag significantly behind the cross-attention fused architecture (0.618). Neither acoustic graphs alone nor text semantics alone suffice; musical context requires joint topological and linguistic grounding.
- **Continuous Emotion Accuracy:** The cross-attention model reduces the Mean Absolute Error (MAE) for valence and arousal prediction from 1.250 down to **0.170**, indicating that cross-modal attention establishes tight affective calibration.
- **Cross-Modal Zero-Shot Retrieval:** The Task 4 contrastive model establishes a Recall@5 of 0.382, demonstrating that the InfoNCE dual-encoder successfully maps structural music graphs into the semantic subspace of descriptive natural-language queries.

5.2 Ablation Study

To evaluate the individual contributions of our architectural design decisions, we conducted comprehensive ablation experiments across four distinct fusion variations:

The ablation results in Table 3 provide critical insights:

1. **Early Concatenation vs. Cross-Attention:** Simply concatenating the global graph embedding g with the $[CLS]$ token yields 0.563 Macro-F1. In contrast, multi-

Table 3: Ablation analysis of multi-modal fusion mechanisms.

Ablation Variant	Macro-F1	AUC-PR	Δ Macro-F1
(1) GNN-only (Structural Audio)	0.521	0.473	-0.097
(2) BERT-only (Lyrical Text)	0.485	0.442	-0.133
(3) Early Concatenation ($[g \parallel \text{CLS}]$)	0.563	0.510	-0.055
(4) Cross-Attention Multi-Modal Fusion	0.618	0.554	Reference




Figure 2: t-SNE Latent Space Visualization of Learned Multi-Modal Representations. The 2D manifold projection demonstrates clear cluster separation across distinct musical genres and mood profiles. Classical and Jazz occupy coherent, well-defined manifolds in the upper projection space, while high-energy Metal and Rock form tight clusters driven by harmonic distortion and rhythmic topology. The smooth transitional boundaries between Pop and Disco illustrate the model’s nuanced contextual understanding.

head cross-attention yields 0.618 (+9.8% improvement). Cross-attention allows the acoustic structure to attend dynamically to individual salient descriptive words (e.g., “distorted”, “melancholic”) rather than relying on a single pooled text vector.

2. Acoustic Dominance in Structural Ambiguity: Comparing GNN-only (0.521) to BERT-only (0.485) indicates that structural audio signals provide a stronger standalone inductive bias for music understanding than text metadata alone.

5.3 Latent Space Topology and Visual Analysis

To investigate the geometric properties of the multi-modal embedding space, we performed t-Distributed Stochastic Neighbor Embedding (t-SNE) on the latent vectors $z \in \mathbb{R}^{2d}$ generated for the test corpus. As illustrated in Figure 2:

- **Manifold Cohesion:** Acoustically and harmonically related genres form distinct topological clusters. Classical compositions (characterized by sparse acoustic similarity edges and prominent chroma variances) are cleanly segregated from rhythmic genres.
- **Semantic-Structural Alignment:** High-energy genres (Metal, Rock) cluster together with high internal cohesion, reflecting both timbral similarity in node features and heavy temporal edge density.

Table 4: Qualitative case studies showing ground truth annotations, model predictions, cross-attention focus tokens, and top-1 contrastive retrieval captions.

Track Example	Ground Truth	Predicted Tags	Top Cross-Attention Tokens	Top Contrastive Retrieved Caption
Track 01 (Rock)	Rock, Energetic ($v = 0.45, a = 0.82$)	Rock, Heavy, Guitar ($\hat{v} = 0.41, \hat{a} = 0.79$)	“distorted”, “heavy”, “driving”, “rhythm”	“An energetic rock song with heavy distorted electric guitars and aggressive drumming.”
Track 02 (Jazz)	Jazz, Smooth ($v = 0.62, a = 0.31$)	Jazz, Relaxing, Brass ($\hat{v} = 0.60, \hat{a} = 0.35$)	“saxophone”, “mellow”, “improvisation”	“Smooth jazz piece featuring mellow saxophone melodies and subtle upright bass accompaniment.”
Track 03 (Classical)	Classical, Somber ($v = -0.38, a = 0.22$)	Classical, Strings ($\hat{v} = -0.34, \hat{a} = 0.25$)	“orchestral”, “sorrowful”, “tempo”	“Melancholic orchestral composition with slow expressive violin phrases and minimal percussion.”

- **Boundary Nuance:** Genres sharing overlapping stylistic heritage (such as Disco and Pop, or Blues and Jazz) exhibit continuous transitional manifolds rather than artificial orthogonal separation, verifying that the model avoids over-clustering.

5.4 Qualitative Case Studies

Table 4 examines qualitative predictions on representative test tracks, demonstrating the interpretive power of cross-modal attention.

6 Discussion, Limitations, & Future Work

While our framework demonstrates substantial empirical gains, several design trade-offs warrant discussion:

- **Fixed Window Granularity:** Segmenting audio into rigid 5-second windows may bisect rapid musical phrases or sustained cadence holds. Future work could incorporate onset-detection algorithms to generate adaptive beat-synchronous segment boundaries.
- **Heuristic Edge Thresholding:** The metric cosine similarity threshold ($\tau = 0.65$) is currently a static hyperparameter. Introducing differentiable graph structure learning (e.g., Gumbel-Softmax edge generation) could allow the network to dynamically discover non-obvious musical connections end-to-end.
- **Full Transformer Fine-Tuning:** Due to memory constraints on commodity hardware, DistilBERT weights were frozen during initial passes. Unfreezing upper transformer layers with low-rank adaptation (LoRA) presents an exciting avenue for further accuracy gains.

7 Broader Impact and Ethical Considerations

Automated musical context understanding carries direct practical applications in streaming recommendation, adaptive music therapy, and accessible multi-modal cataloguing. However, researchers must remain mindful of training dataset biases: Western tonal harmony frameworks dominate standard MIR corpora (GTZAN, MagnaTagATune), potentially penalizing non-Western musical traditions (such as Indian Ragas or Arab Maqam) that operate under alternative tuning systems and cyclical rhythmic concepts. Ethical deployment requires proactive curation of culturally diverse musical graphs.

8 Conclusion

In this paper, we proposed a hybrid structural-semantic architecture for music context understanding. By modeling audio as a Music Structure Graph of acoustic segments connected via temporal adjacency and cosine similarity, our GraphSAGE encoder captures non-local musical topologies that traditional planar CNNs overlook. Coupled with DistilBERT via cross-attention, our model establishes new performance benchmarks: a Macro-F1 of 0.618 (+50% relative over CNNs), an MAE of 0.170 on continuous emotional valence/arousal, and a zero-shot retrieval Recall@5 of 0.382. These results confirm that fusing relational graph message passing with contextual language modeling unlocks a more profound, nuanced comprehension of musical context.

References

- [1] K. Choi, G. Fazekas, and M. Sandler. Convolutional recurrent neural networks for music classification. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 2392–2396, 2017.
- [2] W. Hamilton, Z. Ying, and J. Leskovec. Inductive representation learning on large graphs. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 30, pages 1024–1034, 2017.
- [3] T. N. Kipf and M. Welling. Semi-supervised classification with graph convolutional networks. In *International Conference on Learning Representations (ICLR)*, 2017.
- [4] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 30, pages 5998–6008, 2017.
- [5] V. Sanh, L. Debut, J. Chaumond, and T. Wolf. DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.
- [6] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, et al. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning (ICML)*, pages 8748–8763, 2021.
- [7] A. v. d. Oord, Y. Li, and O. Vinyals. Representation learning with contrastive predictive coding. *arXiv preprint arXiv:1807.03748*, 2018.
- [8] A. Agostinelli, T. I. Denk, Z. Borsos, J. Engel, M. Verzett, A. Caillon, Q. Huang, A. Jansen, A. Roberts, M. Tagliasacchi, et al. MusicLM: Generating music from text. *arXiv preprint arXiv:2301.11325*, 2023.
- [9] Q. Huang, A. Jansen, J. Lee, R. Ganti, J. Y. Chen, and E. Humphrey. MuLan: A joint embedding of music audio and natural language. In *International Society for Music Information Retrieval (ISMIR)*, 2022.
- [10] G. Tzanetakis and P. Cook. Musical genre classification of audio signals. *IEEE Transactions on Speech and Audio Processing*, 10(5):293–302, 2002.
- [11] B. McFee, C. Raffel, D. Liang, D. P. Ellis, M. McVicar, E. Battenberg, and O. Nieto. librosa: Audio and music signal analysis in python. In *Proceedings of the 14th Python in Science Conference*, pages 18–25, 2015.
- [12] M. Fey and J. E. Lenssen. Fast graph representation learning with PyTorch Geometric. In *ICLR Workshop on Representation Learning on Graphs and Manifolds*, 2019.
- [13] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, et al. PyTorch: An imperative style, high-performance deep learning library. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 32, pages 8024–8035, 2019.
- [14] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *North American Chapter of the Association for Computational Linguistics (NAACL)*, 2019.
- [15] P. Veličković, G. Cucurull, A. Casanova, A. Romero, P. Liò, and Y. Bengio. Graph Attention Networks. In *International Conference on Learning Representations (ICLR)*, 2018.
- [16] J. Serrà, Á. Corral, M. Boguñá, M. Haro, and J. L. Arcos. Measuring the evolution of contemporary western popular music. *Scientific Reports*, 2(1):521, 2012.